This is going to be a comprehensive interview lesson, Let me build it properly — deep explanations, visual diagrams and real-world points.

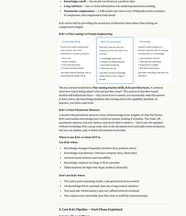

1. What is R.E.S.?

Restful Application Interface is a technique that gives a LLM access to external tools and data sources through a standardized HTTP interface, instead of hardcoding a plugin or processing context before generating an answer.

Think of it this way: an LLM without R.E.S. is like a doctor who graduated medical school 20 years ago and has never seen a patient since. It knows everything it learned, but it can't look up the latest drug interaction before prescribing you something.

Why do LLMs hallucinate?

Because they are trained on **static knowledge**. They don't "know" facts — they've learned statistical correlations between tokens from billions of examples. When asked something, they generate the most plausible continuation. If the correct answer wasn't in the training data, or the model's knowledge is outdated, it will still confidently produce the most plausible-sounding wrong answer.

The three causes:

- **Knowledge cutoff** — the model knows facts as of a point in time
- **Long tail bias** — rare facts are underrepresented in training
- **Parametric compression** — a 70B model can't store every Wikipedia article, it lossily compresses patterns into weights

LLM context limits exacerbate this by actual hard architectural limits rather than relying on memory alone.

RAG vs Tool Calling vs Fine-Tuning

Fine-tuning	RAG	Tool calling
Bake knowledge into model weights. Slow to update, expensive to retrain, opaque to inspect.	Retrieve relevant text from a vector database and inject into context before generation. Fast to update, cheap, transparent.	Model calls a function at runtime to fetch live data or take an action. Real-time, structured, composable.
Best for: tone, style, domain vocabulary, teaching the model a new task format.	Best for: large knowledge bases, frequently changing facts, citations needed.	Best for: live data, actions, anything that changes between requests.
Can't add new knowledge without retraining. Can't cite sources. Can't update without a new deploy.	Adds latency (retrieval step). Quality depends on chunking and embedding choices. Can't reason across many documents.	Model must know when to call and how to interpret results. Adds latency per tool call. Failure modes multiply.

The key second-order trade-off: **fine-tuning makes the model behave differently, RAG gives it more to work with, tool calling lets it act on the world.** These are genuinely different problems and you'll combine all three in production systems.

R.E.S. vs Function Calling

Function calling is a vendor-specific API feature (OpenAI, Anthropic, Google). R.E.S. is a conceptual layer that sits on top of function calling. It adds:

- Standardized tool discovery (the model doesn't need to know in advance what tools exist)
- Capability negotiation (tools can declare what they support)
- Composability (multiple servers can be combined into one interface)
- Vendor independence (the same tool definition works across providers)

Function calling is the mechanism. R.E.S. is the architecture that makes it composable and portable.

When to use R.E.S. vs direct API integration

- Prototyping or single-provider integrations: direct function calling is fine, simpler to ship
- Multi-provider, multi-tool, or shared tool ecosystem: R.E.S. pays for itself in reusability
- Enterprise with many internal tools: R.E.S. lets teams publish tools once and any agent can use them

Why Anthropic created it

Anthropic's Claude models were trained to follow the R.E.S. protocol natively. This means Claude is better at deciding when to call tools, parsing tool results, and recovering from tool errors than models trained only on function calling syntax.

2. Core R.E.S. Primitives — Each Phase Explained

Initializing Phase — Client Side

This runs once when a user begins a conversation — or incrementally as new data arrives.

Tool Discovery — the client asks the server "what tools do you have?" The response includes JSON Schema definitions for each tool, including its parameters. The model will ultimately decide which tool to call based on these descriptions — they must be clear, accurate, and specific. Vague descriptions like "handle user data" will cause the model to call the wrong tool or hallucinate parameters.

Initializing Phase — Server Side

Capability Negotiation — the server declares what it supports. Common capabilities include:

- Knowledge retrieval (semantic search, keyword search)
- Knowledge ingestion (internal company docs, client data)
- Tool execution (API calls, database queries)
- Knowledge creation (writing back to a knowledge base)
- Authentication (per-tenant keys, OAuth, mTLS)

Prompt Phase — Client Side

- The user's prompt reaching Claude, with generation from scratch
- Or continuing from a checkpoint (just one step, not the entire conversation)
- The system prompt and tool definitions are sent with every request
- The context window fills with tokens (just like a chat with no memory at all)

Prompt Phase — Server Side

Tool Search — the server receives the user's prompt and runs a search against its knowledge base. This is the retrieval step. The quality of this step matters more than anything else in the system. Bad retrieval means the model gets irrelevant tokens and wastes them. Good retrieval means the model gets exactly what it needs to reason and act.

Tool Ranking — if many tools matched, the server can rerank them by relevance before returning them to the model. Cheap models can do this well.

Tool Injection — the server returns the matching tool definitions. The model now knows what it can call and how to call it.

Execution Phase — Client Side

Tool Call — Claude decides to call a tool and emits a structured request, for example:

```
search_knowledge_base("Q3 revenue", filters={"year": 2024})
```

Tool Execution — the client runs the actual function against the server. Latency here depends entirely on what the tool does — a database query is fast, a web scrape is slow, a human approval is infinite.

Execution Phase — Server Side

Result Formatting — the server formats the raw output into something the model can read efficiently. Markdown tables are more token-efficient than JSON for structured data. Truncate long outputs with a clear indicator that more exists.

Context Injection — the formatted result goes into the conversation as a tool result message. The model now sees it alongside everything else in the context window.

Continuation — Client Side

Synthesis — Claude reasons over the tool results and generates a response, or calls another tool if it needs more information. Multi-step reasoning happens here.

Stream Delivery — the response streams back to the user token by token.

Continuation — Server Side

Logging — every tool call, latency, and outcome is recorded for observability. Without this, you can't debug why the model made a bad decision three weeks later.

Cache Warming — if the same queries are common, pre-fetch results into a cache before the next request arrives.

3. R.E.S. Variants — The Evolution

Each variant solves a specific pain point of plain R.E.S. Know how to talk about them.

Plain R.E.S.	Stateless request-response. Simple, cheap, cacheable. Can't remember anything between calls.
Streaming R.E.S.	Adds SSE for progress updates and partial results. Better UX for long-running operations.
Bidirectional R.E.S.	Persistent connection, server can push notifications. Enables real-time collaboration and monitoring.
Agentic R.E.S.	Adds planning, memory, and multi-step reasoning. Server manages state across calls. Closer to a full agent than a tool interface.
Multi-Agent R.E.S.	Multiple specialized agents behind one interface. Orchestrator routes to specialists. Enterprise-grade architecture.

4. Core Components — What Each Does Once

This is where interviewers test whether you can implement, not just describe.

Document Loader — Load raw text from PDFs, HTML, Markdown, databases, APIs. Preserve structure (headings, sections, tables) where possible. Output: `RawDocument(text, metadata)`

Text Splitter / Chunker — Split documents into chunks small enough to embed and retrieve precisely. Chunk too large: retrieval returns noisy context. Chunk too small: lose surrounding context. Typical: 512–1024 tokens with 10–20% overlap. Output: `List[Chunk(text, metadata, parent_id)]`

Embedding Model — Convert chunk text into a dense vector. Must be the same model at index time and query time. Output: `List[float]`

Vector Store — Stores chunk vectors + metadata. Supports approximate nearest neighbor search. Examples: Pinecone, Weaviate, Qdrant, pgvector, Chroma. Output: `List[SearchResult(chunk_id, score)]`

Retriever — Orchestrates the search: embeds the query, calls the vector store, applies filters, optionally runs keyword search (BM25) alongside vector search, reranks results. Output: `List[Chunk]`

Reranker — Cross-encoder that scores query + chunk together for more accurate relevance than embedding similarity alone. Slow but high-quality. Output: `List[Chunk(..., rerank_score)]`

Context Assembler — Formats retrieved chunks into the prompt, respecting token budgets. Decides what to include when results exceed the context window. Output: `str`

LLM Generator — Takes assembled context + query, generates grounded answer. Can generate citations referencing specific chunks. Output: `Answer(text, citations, confidence)`

Cache Layer — Stores embeddings, retrieval results, and generated answers. Reduces latency and cost. Critical for production. Output: `Optional[CachedResult]`

Ingestion Pipeline — Orchestrates the full flow: load → split → embed → store. Handles batching, rate limits, retries, and incremental updates. Output: `IngestionReport(docs_processed, chunks_created, errors)`

5. Chunking Strategies — When Each Wins

This is where interviewers test whether you can implement, not just describe.

Strategy	When it wins	Trade-off
Fixed-size	Quick baseline, homogeneous text	Splits mid-sentence, ignores structure
Recursive	Preserves paragraphs and sentences when possible	Still ignores semantic boundaries
Semantic	Topic shifts matter, heterogeneous documents	Slow, requires embedding during ingestion
Parent-child	Need precise retrieval but broad context for reasoning	Storage overhead, more complex retrieval logic
Late chunking	Long-context embedding models available	Requires specific model support

Fixed-size chunking splits every N characters regardless of structure. Fast and predictable, but it will cut a sentence in half and lose the thread entirely.

Recursive chunking tries paragraphs first, then sentences, then words, then characters, until chunks fit the target size. Preserves structure when possible. Most production systems start here.

Semantic chunking embeds sentences and splits where the embedding similarity drops sharply between adjacent sentences. Slow at ingestion time but produces chunks that are topically coherent.

Parent-child chunking stores small chunks for precise retrieval but returns their parent paragraph or section for context. Solves the tension between "I need to find the right thing" and "I need enough context to reason about it."

Late chunking embeds the full document first, then splits the resulting token embeddings. Preserves cross-chunk context that other methods lose. Requires a long-context embedding model.

Interview Question Bank — Domain 2

- Walk me through what happens when a user asks a question in your system.
- How do you decide chunk size for your system?
- When would you choose BM25 over vector search?
- What is a reranker and when is it worth the latency cost?
- How do you handle stale data in your knowledge base?
- Why does the same embedding model need to be used at index time and query time?
- How would you design a system that cites its sources?
- Explain BM25 in 3 bullets — what problem does it solve?

Quick Mental Models to Internalize

Before you start coding, these three ideas will save you from most mistakes:

Retrieval is a ranking problem, not a lookup problem. You're not looking up a fact, you're ranking thousands of candidates by relevance and hoping the right one is on top.

Embedding + LLM is probabilistic by nature. Just like LLM outputs aren't perfect, neither is retrieval. Design for graceful degradation when the system doesn't find what it needs.

Chunking is 80% of quality. If you fix one thing in a broken retrieval system, it's almost always chunking. Bad chunks can't be rescued by better embeddings or better prompts. Diagnose, then fix, the chunking first.

Tier 1 LLM capabilities matrix (Claude, GPT-4, Gemini): Check the mathematical ceiling of your model family before you start. Some models can't do multi-step reasoning regardless of how well you write the prompt. That's the right trade-off for simple tasks, but it will fail on complex ones.

Re-rankers — Second-stage. The 10 results you retrieved by similarity search are fast but noisy. Re-rankers read each candidate against the query and reorder by true relevance. Adds latency per search, but the top-5 they hand to the LLM is significantly better than the top-5 they'd hand without one.

Then it asks: "Alright, I'm ready. Start with Domain 1 (Chunking Strategies) or do you want to jump to Domain 2 (Retrieval Systems)?"

Let's go deep. I'll build this as a layered teaching system — concept first, then mechanics, then the interview questions you'll actually get.

Why chunking is the most underrated problem in R.E.S.

People obsess over which LLM to use, which vector DB to pick, how to tune the prompt. Meanwhile they chunk documents by splitting every 512 characters regardless of whether that cuts a sentence in half or separates a question from its answer.

Chunking is the foundation everything else sits on. Bad chunks can't be rescued by better embeddings, better prompts, or better models. The LLM can't reason about information it never receives. Garbage in, garbage out, but the garbage is invisible until you trace a wrong answer all the way back to ingestion.

1. Fixed-Size Chunking

The simplest strategy. Cut every N characters, ignore structure, add overlap to soften the damage.

Character-based splitting counts raw characters. Fast, deterministic, predictable. Breaks mid-word, mid-sentence, mid-paragraph. Acceptable for logs or machine-generated text. Terrible for prose.

Token-based splitting counts tokens instead of characters. Aligns better with how the model actually sees text. Still breaks structure, but at least the chunk boundary matches what the model cares about when it reads the input.

Overlap (typically 10–20%) means each chunk shares some tokens with its neighbor. This softens the damage when a sentence or idea gets split across two chunks — the overlap ensures enough context survives on both sides to be useful.

Pros: Simple to implement, fast to run, deterministic, predictable memory usage.

Cons: Ignores structure entirely, breaks mid-sentence, wastes tokens on overlap, performs poorly on anything with meaningful structure.

Use when: Logs, machine-generated text, as a baseline you measure everything else against.

2. Recursive Splitting — The Production Default

Try to split on paragraph boundaries first. If the resulting chunks still exceed the target size, recurse into sentences, then words, then characters. Stop as soon as everything fits within the target size.

This respects structure by default. Paragraphs stay intact when they can, sentences when they can't, words only when forced.

```
Recursive(text, target=512)
  ├─ split on "\n\n"  → [para1, para2, para3]
  │   para1 fits? ✓
  │   para2 fits? ✗ → recurse into para2
  │     split on "\n" → [sentence1, sentence2]
  │     sentence1 fits? ✓
  │     sentence2 fits? ✗ → recurse into sentence2
  │       split on " " → [word1, word2, word3, ...]
  │       ...
  ├─ para3 fits? ✓
```

Recursive Separator Hierarchy — separators tried in order

Stop as soon as everything fits within the target size.

Pros: Respects structure when possible, adapts to document type, preserves more meaning than fixed-size.

Cons: Still ignores semantic boundaries, produces variable chunk sizes, harder to reason about than fixed-size.

Use when: Prose documents, articles, documentation, anything with paragraphs and sections. This is the default for most RAG systems — it's simple enough to ship in an afternoon and good enough for most corpora.

3. Semantic Chunking — Splitting When Meaning Changes

Embed each sentence, compute cosine similarity between adjacent sentences, and split where similarity drops sharply. Sentences that are about the same thing stay together. Sentences that are about different things get separated.

```
sentences = ["Revenue grew 3% in Q3.", "Gaming revenue was the strongest.",
             "The CEO retired in 2024.", "Her successor is a former CFO."]

embeddings = [embed(s) for s in sentences]

similarity(1, 2) = 0.82   # revenue → gaming: same topic
similarity(2, 3) = 0.31   # gaming → CEO: topic shift
similarity(3, 4) = 0.79   # CEO → successor: same topic

→ chunk 1: "Revenue grew 3% in Q3. Gaming revenue was the strongest."
→ chunk 2: "The CEO retired in 2024. Her successor is a former CFO."
```

Trigger between the three threshold standard deviations below to mark a boundary. Higher values produce fewer, larger chunks. Lower values produce more, smaller chunks.

Pros: Preserves semantic boundaries, adapts to document structure, produces coherent chunks.

Cons: Requires embedding every sentence at ingestion time, slow and expensive at scale, sensitive to embedding model quality, produces variable chunk sizes.

Use when: Long-form prose where topic shifts matter, when you have budget for ingestion compute, when retrieval quality is worth the cost.

4. Parent-Child Chunking — Decoupling Retrieval from Generation

The tension: small chunks retrieve precisely but lose context. Large chunks give context but retrieve noisily. Parent-child resolves this by storing both.

Index small chunks for precise retrieval. When a chunk matches, return its parent paragraph or section to the LLM for reasoning.

```
Document
  ├─ Section A
  │   ├─ Paragraph 1  → embed → vector_1
  │   ├─ Paragraph 2  → embed → vector_2
  │   ├─ Paragraph 3  → embed → vector_3
  │   ├─ Paragraph 4  → embed → vector_4
  │   ├─ Paragraph 5  → embed → vector_5
  │   ├─ Paragraph
```